

# A screen space GPGPU surface LIC algorithm for parallel interactive exploration of large composite datasets

Astronum 2014

B. Loring LBNL, H. Karimabadi and V. Roytershteyn UCSD, W. Daughton LANL

March 16, 2015

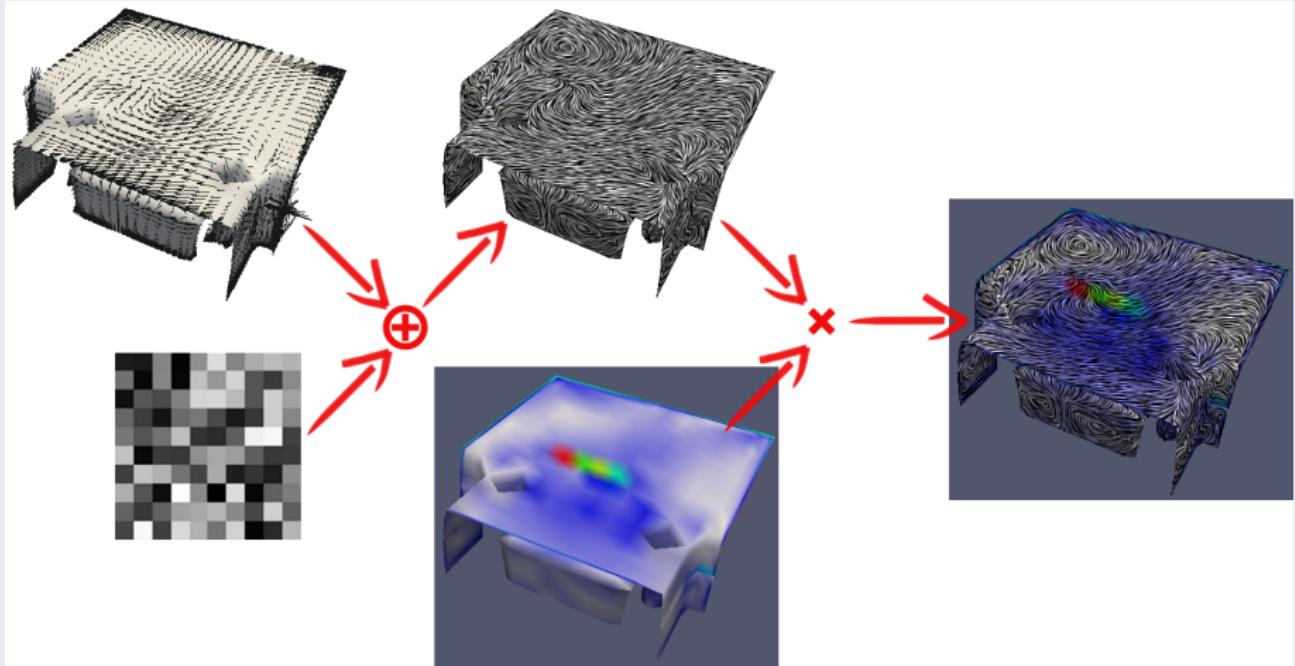
## Goal of the work

We set out to design a surface LIC algorithm that:

1. handles all common datasets, including AMR, multiblock etc...
2. is parallel, can process large simulation data
3. deployable in existing vis tools
4. delivers interactive performance
5. is flexible enough to work well on wide range of data
6. exposes tuning parameters to the user

# Introduction

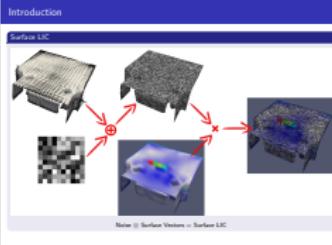
## Surface LIC



Noise  $\otimes$  Surface Vectors = Surface LIC

# A screen space GPGPU surface LIC algorithm for parallel interactive exploration of large composite datasets

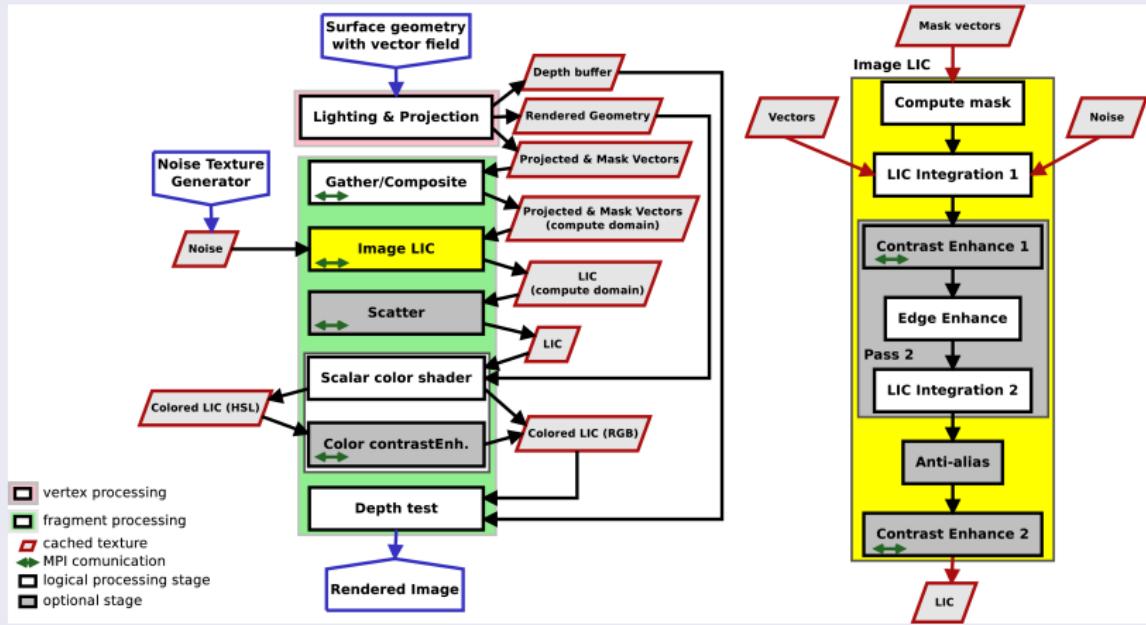
## └ Introduction



- \* global view of the vector field, easily identify features
- \* computationally costly, GPGPU implementation to make it interactive
- \* working in screen space introduces some interesting complexity and challenges
- \* especially for the parallelization, if you expect it to work in existing parallel vis tools

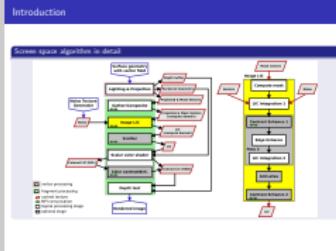
# Introduction

## Screen space algorithm in detail



# A screen space GPGPU surface LIC algorithm for parallel interactive exploration of large composite datasets

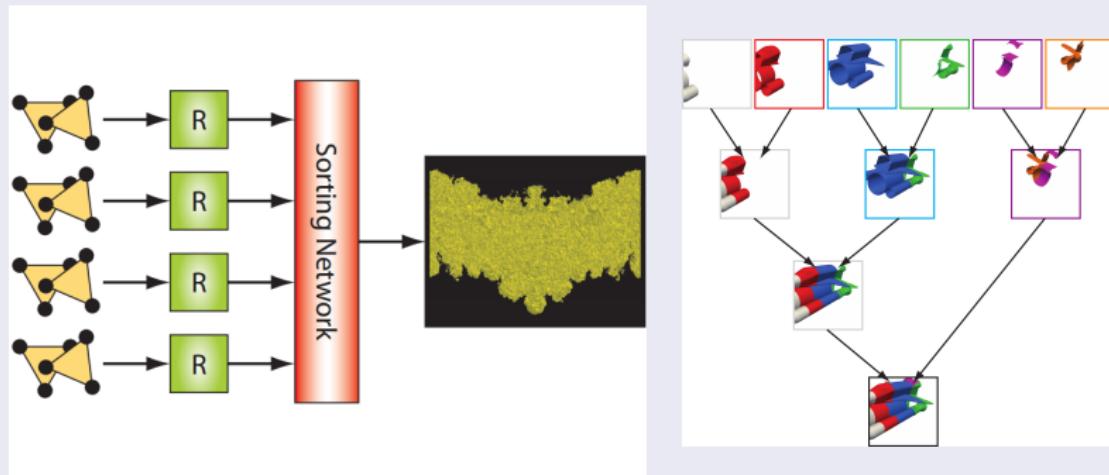
## └ Introduction



- \* lighting and projection shader
- \* screen space after this
- \* compositing stages for parallel
- \* image LIC is broken out on the right
- \* 2 pass image LIC, first pass convolves noise, second pass convolves first pass after image processing filters
- \* some of the image processing stages require guard pixels for correct parallel operations
- \* some require reduction
- \* shaders for combining LIC and rendered geometry
- \* stages are cached
- \* works with IceT image compositing library used by ParaView and VisIt

# Parallelization

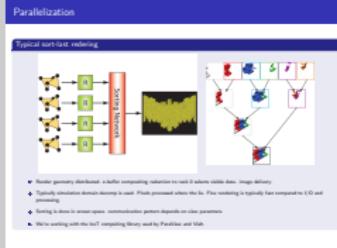
## Typical sort-last rendering



- Render geometry distributed. z-buffer compositing reduction to rank 0 selects visible data. image delivery.
- Typically simulation domain decomp is used. Pixels processed where the lie. Fine rendering is typically fast compared to I/O and processing.
- Sorting is done in screen space. communication pattern depends on view parameters.
- We're working with the IceT compositing library used by ParaView and VisIt.

# A screen space GPGPU surface LIC algorithm for parallel interactive exploration of large composite datasets

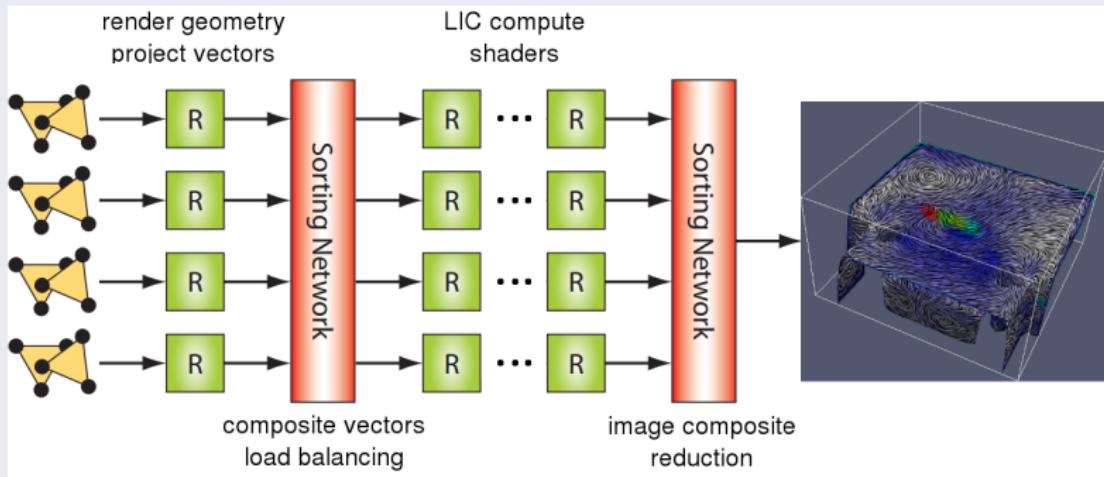
## └ Parallelization



distributed geometry using simulation domain decomp  
screen space domain decomposition is view dependent  
cloesest pixels need to be displayed thus z-buffer compositing

# Parallelization

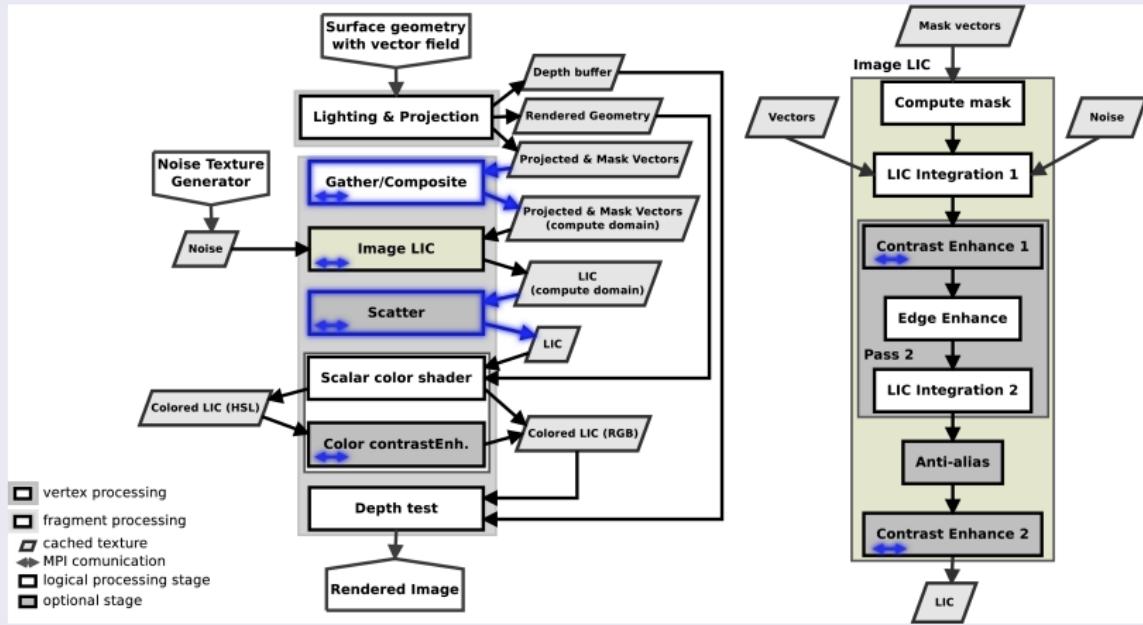
## Sort last with screen space surface LIC



- prior to integration composite vectors where screen space projections overlap
- need “guard pixels” at screen space block boundaries
- shader communication, eg. to get min/max for histogram stretching

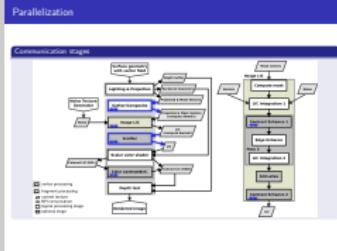
# Parallelization

## Communication stages



# A screen space GPGPU surface LIC algorithm for parallel interactive exploration of large composite datasets

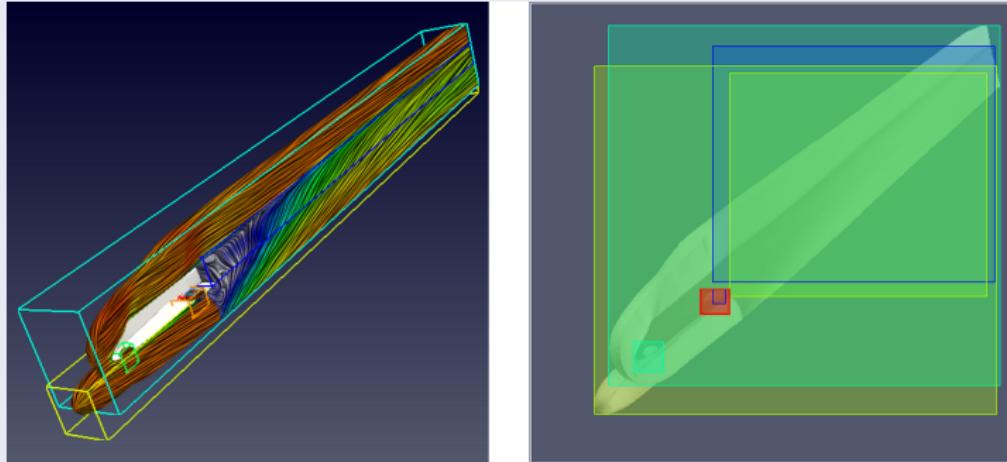
## └ Parallelization



\* compositing + ghost generation done in the gather stage \* image processing stages need min/max \* scatter is a transfer only op, no compositing \* scatter stage is needed \*if\* we do any screen space load balancing because IceT pre-computes the image space bounds \* this could be addressed so that the scatter is not needed, but then we need to composite color buffer so its actually not that bad

# Parallelization

## Working in screen space

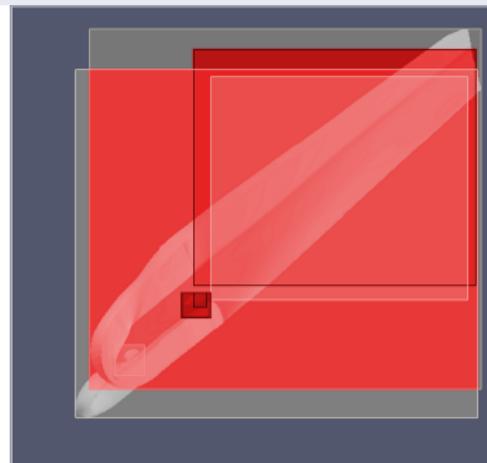
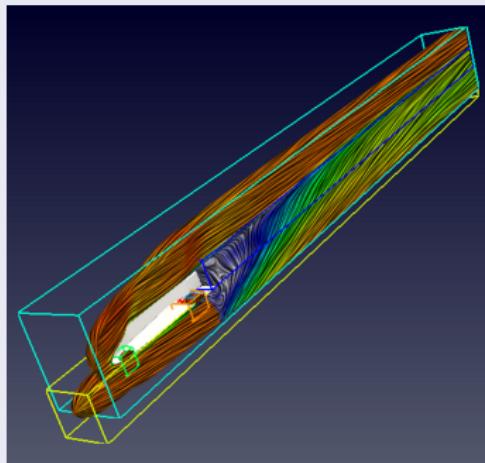


**In-place Composite** Lazy. Process on the existing screen space extents.

- ✓ Simple, works with IceT
- ✓ It's the best if there's no overlap (ie slice)

# Parallelization

## Working in screen space



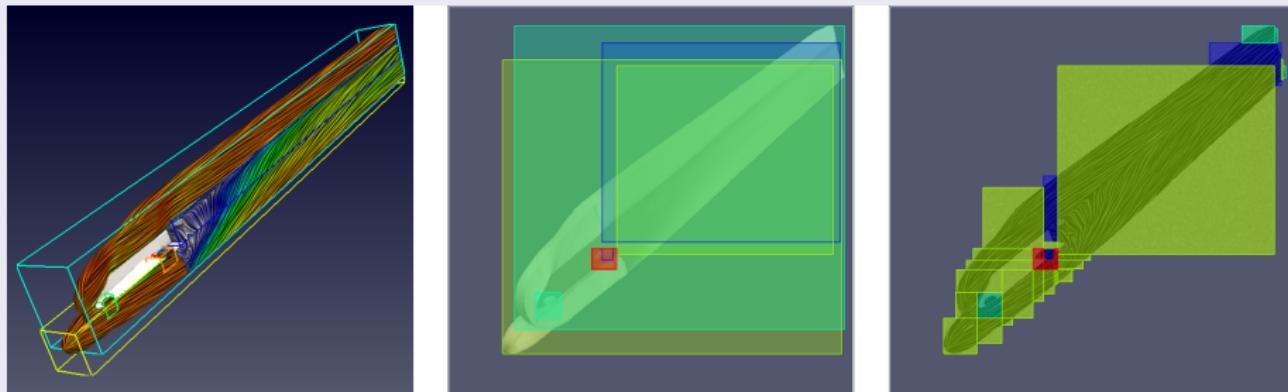
**In-place Composite** Lazy. Process on the existing screen space extents.

- ✓ Simple, works with IceT
- ✓ It's the best if there's no overlap (ie slice)
- ✗ view dependent/not load balanced
- ✗ high communication costs. all process pairs with overlapping regions exchange data
- ✗ redundant computation can limit performance

*screen-space overlap is costly!!*

# Parallelization

## Working in screen space

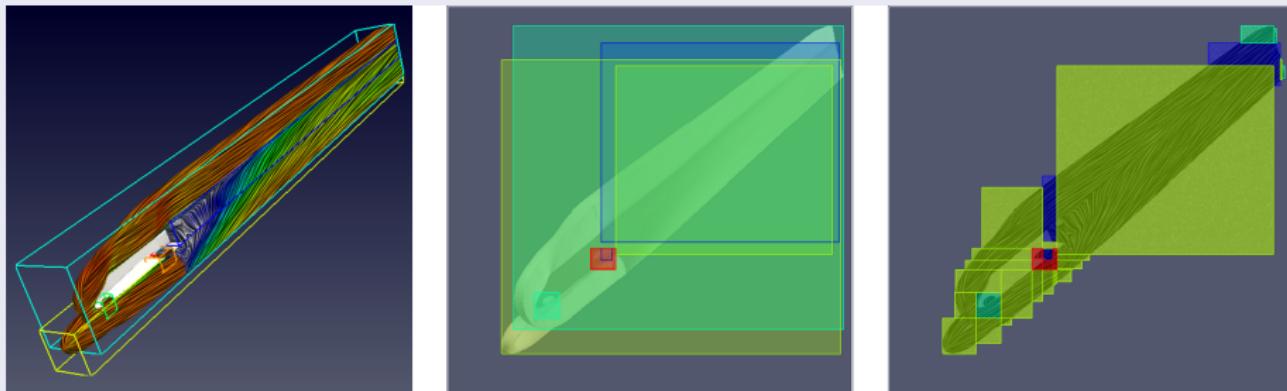


In-place Disjoint Composite *Less lazy. Process on *unique* parts of existing screen space ext.*

- ✓ reduced communication costs by  $\frac{1}{2}$  for each overlapping processor pair
- ✓ computing once per pixel is a big win, esp. w/o GPUs
- ✓ smaller regions further reduce comm and compute costs

# Parallelization

## Working in screen space

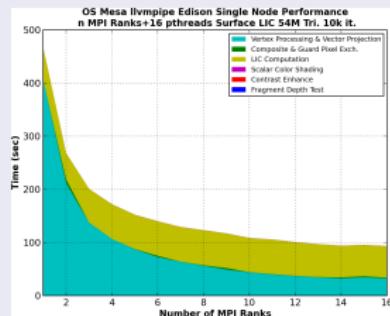
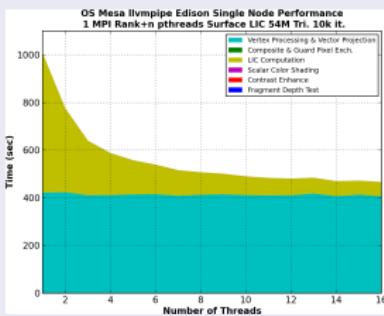
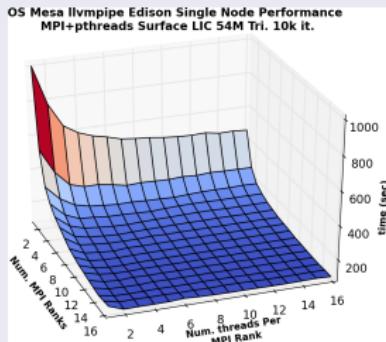


In-place Disjoint Composite *Less lazy. Process on *unique* parts of existing screen space ext.*

- ✓ reduced communication costs by  $\frac{1}{2}$  for each overlapping processor pair ✗ disjointification adds some overhead
- ✓ computing once per pixel is a big win, esp. w/o GPUs ✗ more patches means more guardpixels means more communication
- ✓ smaller regions further reduce comm and compute costs ✗ doesn't solve load balancing issues

# Parallelization

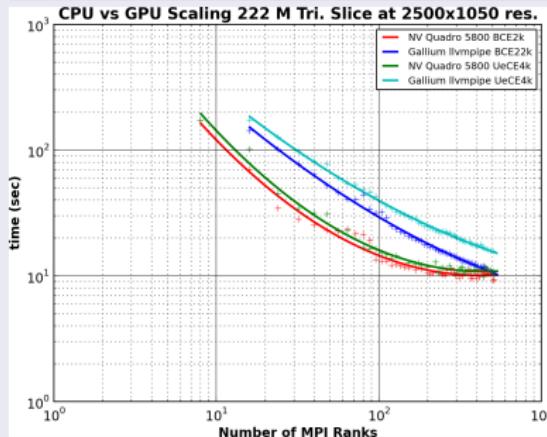
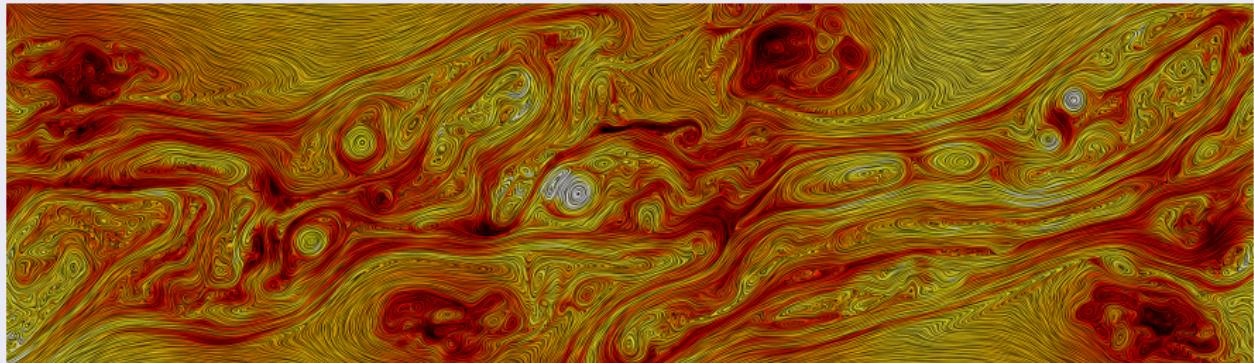
## Single node performance



- w/o GPU
  - Best results when number of threads on the node is equal to the number of hyperthreads.
  - For benchmarks above and below run 1/2 packed, 4 processes per socket, each with 4 threads.
- w GPU
  - Dual GPU nodes
  - Run fully packed 4 processes per GPU

# Parallelization

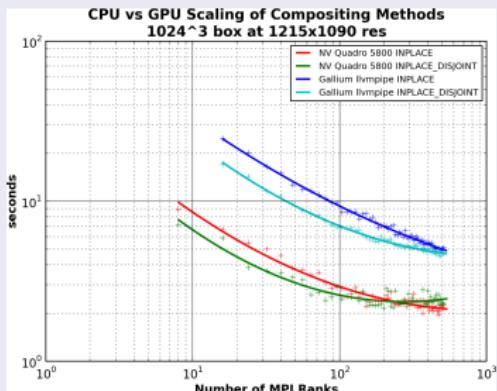
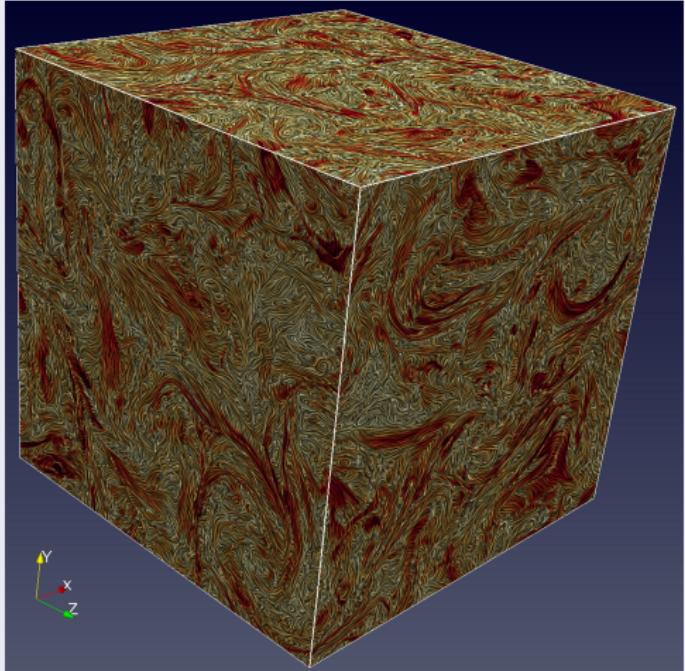
## Performance and Scaling w/wo GPU



- 2D 16384x1890 cell PIC plasma turbulence simulation.  
7TB of data generated 40000 cores in 72 hours on Jaguar.
- 222 Million triangles. 2550x1050 rendering
- GPU faster but CPU scales better
- CPU is really slow!
- GPU needs work to amortize the data transfers over PCI
- more processes → more guard pixels → more comm
- Edison has faster lower latency network hardware

# Parallelization

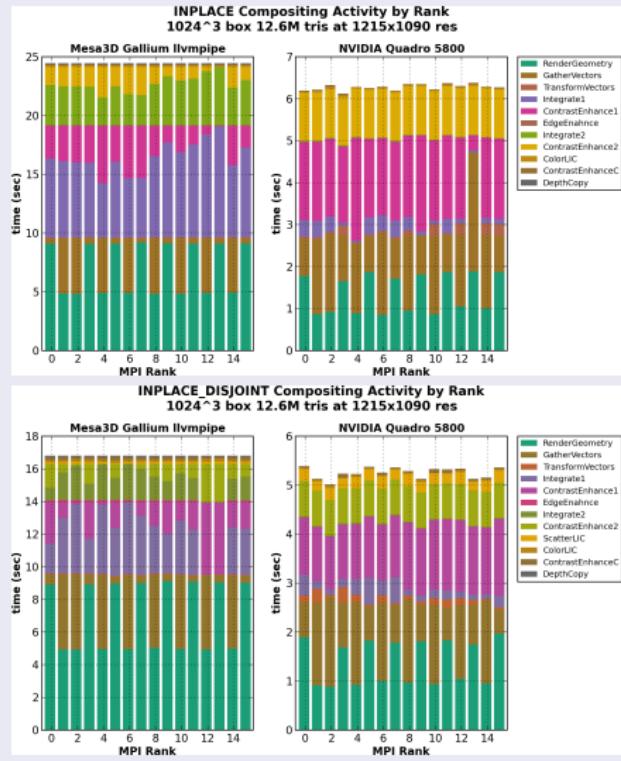
## Performance and Scaling of Compositing Alg w/wo GPU



- 3D 1024x1024 MHD plasma turbulence simulation
- 54 Million triangles. 1215x1090 rendering
- GPU faster but CPU scales better
  - Same as before
  - Disjointification may be resulting more guard pixels and fragmentation

# Parallelization

## Performance and Scaling of Compositing Alg w/wo GPU

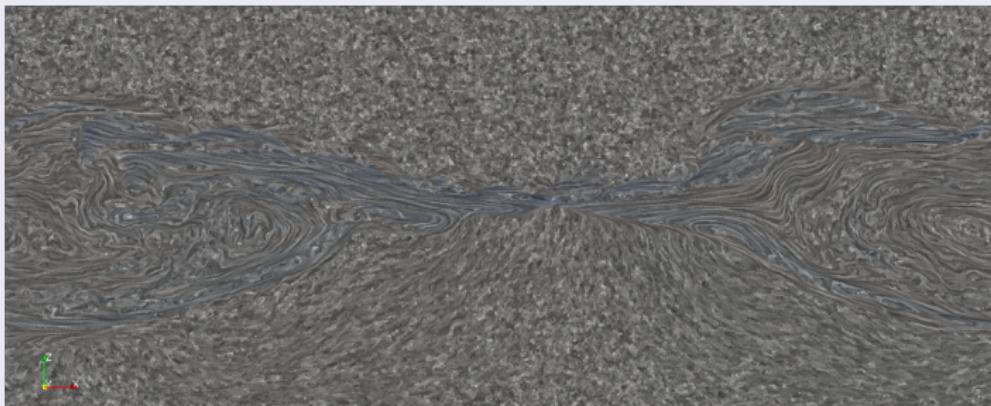


- 16 rank run, slowest run, differences are easier to spot
- first column : w/o GPU , disjointification shaves off time in integration (purple) and communication(brown), scatter stage(yellow second row) adding very little overhead
- second column : w GPU , communication is the bottleneck, gather (brown), contrast enhance (pink, yellow).

# Addressing visual and perceptual issues

## Issues addressed by our work

### Motivation:



### Issues:

- Results are highly dependent on inputs, many inputs, noise, vector, surface, lighting, view, camera, etc
- Convolution of noise naturally reduces contrast and narrows dynamic range
- Combination of LIC with pseudocolored lit geometry often loses subtle variations in pseudocoloring, lighting.

**What properties in LIC/shaders enable best combination with lit pseudocolored geometry?**

# Addressing visual and perceptual issues

First: How is LIC combined with lit pseudocolored geometry?

**Blend:**

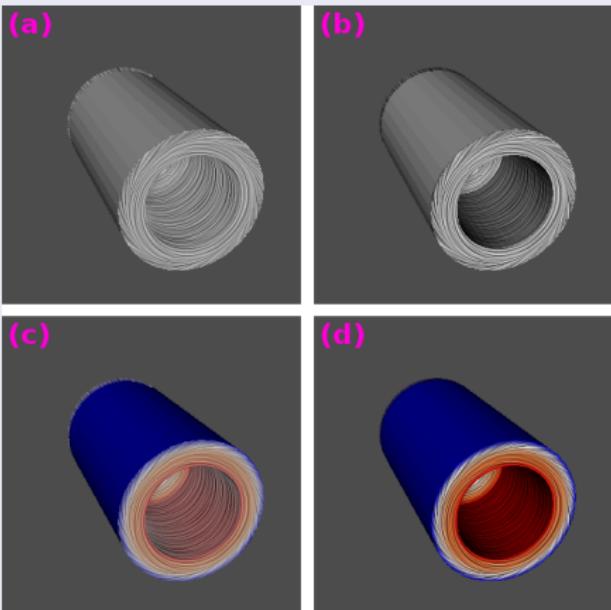
$$c_{ij} = L_{ij} * I + S_{ij} * (1 - I) \quad (1)$$

where the indices  $i, j$  identify a specific fragment,  $c$  is final RGB color,  $L$  is LIC gray scale value,  $S$  is the scalar RGB color, and  $I$  is a constant ranging from 0 to 1.

**Multiply:**

$$c_{ij} = (L_{ij} + f) * S_{ij} \quad (2)$$

where the indices  $i, j$  identify a specific fragment,  $c$  is final RGB color,  $L$  is LIC gray scale intensity,  $S$  is the scalar RGB color, and  $f$  is a biasing parameter, typically 0, that may be used for fine tuning.



**Properties in LIC enable best combination with lit pseudocolored geometry:**

- LIC should have high contrast streaks with lots of pure white and black pixels.
- The combination process needs to preserve high contrast and dynamic range.

# A screen space GPGPU surface LIC algorithm for parallel interactive exploration of large composite datasets

## Addressing visual and perceptual issues

### Addressing visual and perceptual issues

First, how is LIC combined with 3D perspective-correct geometry?

#### Blend

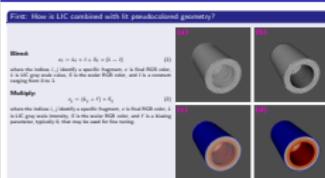
$\alpha_1 = \delta_1 \times t + \delta_2 \times (1-t)$  (1)  
where the index  $\lfloor \cdot \rfloor$  identify a specific fragment,  $t$  is final RGB value,  $\delta_1$  &  $\delta_2$  are the two static RGB const., and  $\alpha_1$  is a constant ranging from 0 to 1.

#### Multiply

$\alpha_1 = (\delta_1 \times t) \times \delta_2$  (2)  
when the index  $\lfloor \cdot \rfloor$  identify a specific fragment,  $t$  is final RGB value,  $\delta_1$  &  $\delta_2$  are the two static RGB const., and  $\alpha_1$  is a constant ranging from 0 to 1.

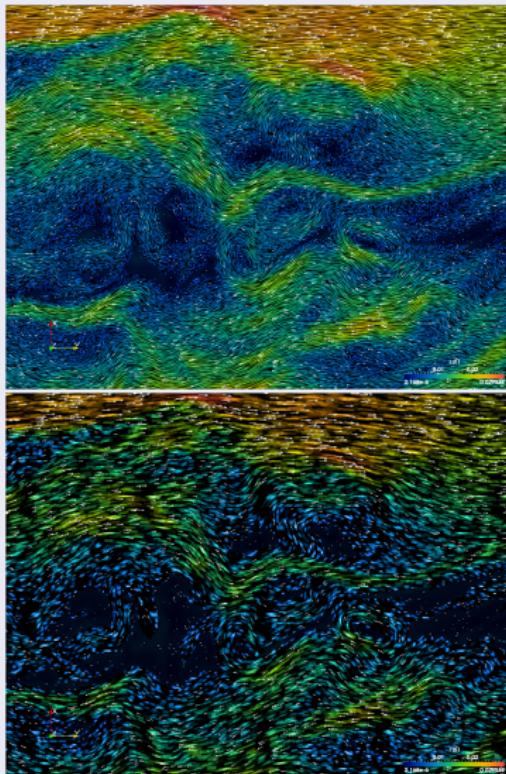
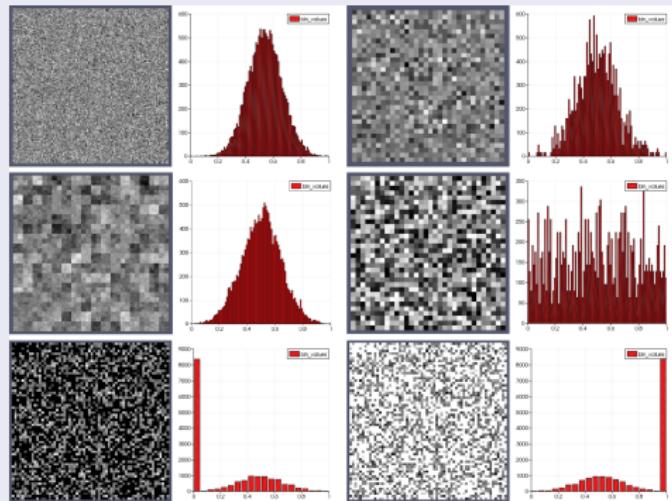
Properties in LIC enable local combinations with 3D perspective-correct geometry:

- LIC algorithm has high contrast when sets of pixels are at ideal position
- The combination process needs to preserve high contrast and dynamic range



# Addressing visual and perceptual issues

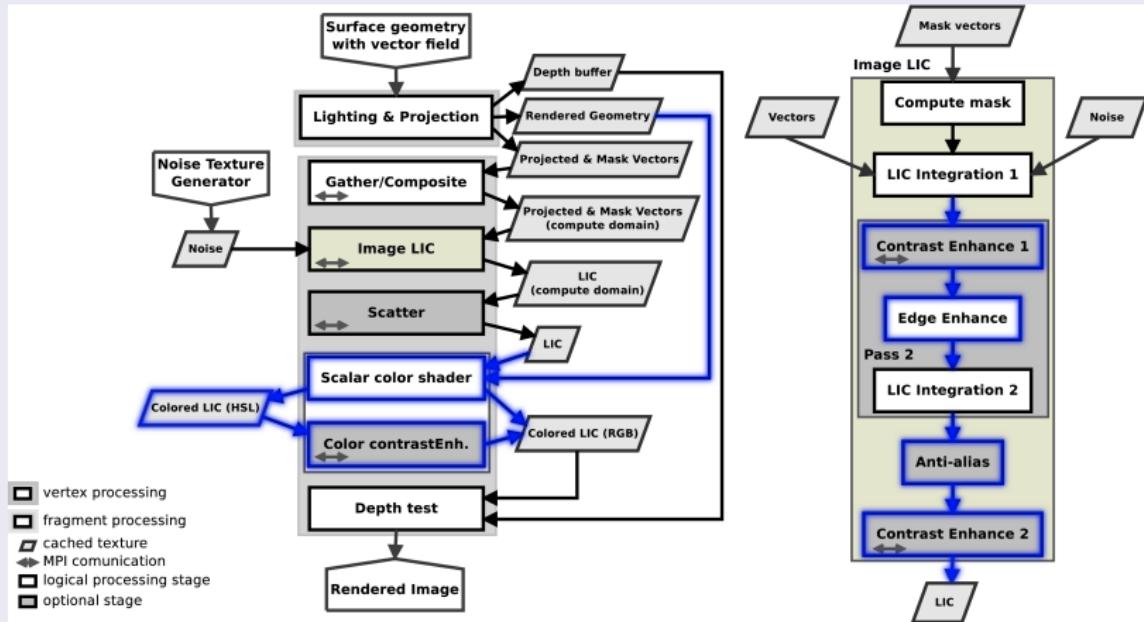
## Noise distribution and properties



- 9 DOF
  - size, distribution, num levels, grain size, min/max value, impulse probability, background color, seed
- better control over streaking. size, width, intensity

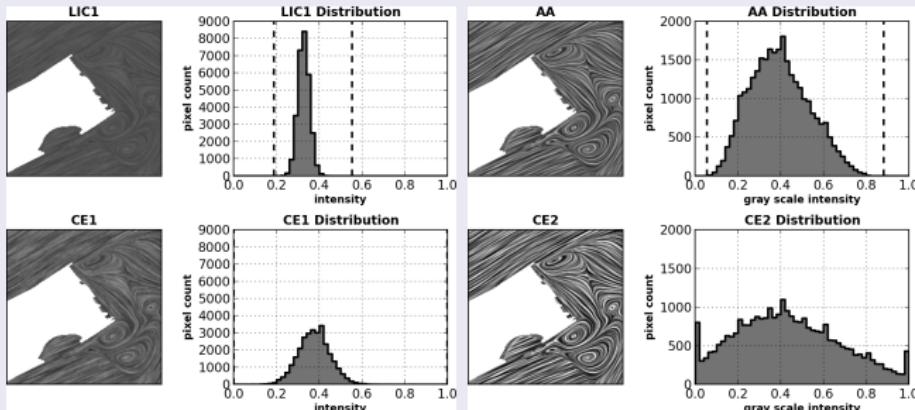
# Introduction

Screen space + OpenGL  $\Rightarrow$  use image processing algorithms!



# Visual and perceptual issues

Histogram Stretching:  
enhances streak contrast and can increase numbers of pure white and black pixels



## Linear re-mapping of intensity values

$$I_{ij} = \frac{I_{ij} - m}{M - m} \quad (3)$$

where, the indices  $i, j$  identify a specific fragment,  $I$  is the fragment's gray scale intensity,  $m$  is the intensity to map to 0,  $M$  is intensity to map to 1.

## Automating the procedure

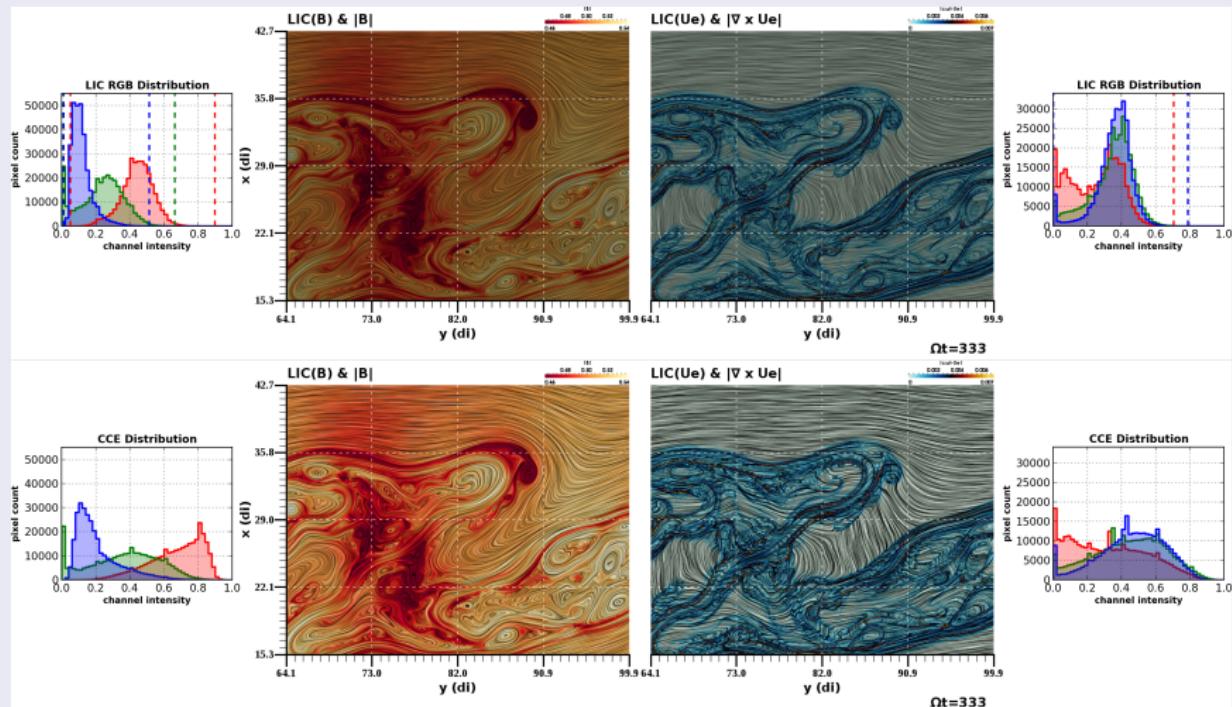
$$m = \min(I) + F_m * (\max(I) - \min(I)) \quad (4)$$

$$M = \max(I) - F_M * (\max(I) - \min(I)) \quad (5)$$

where,  $\min/\max$  are computed on all fragments,  $F_m$  and  $F_M$  are optional adjustment factors.

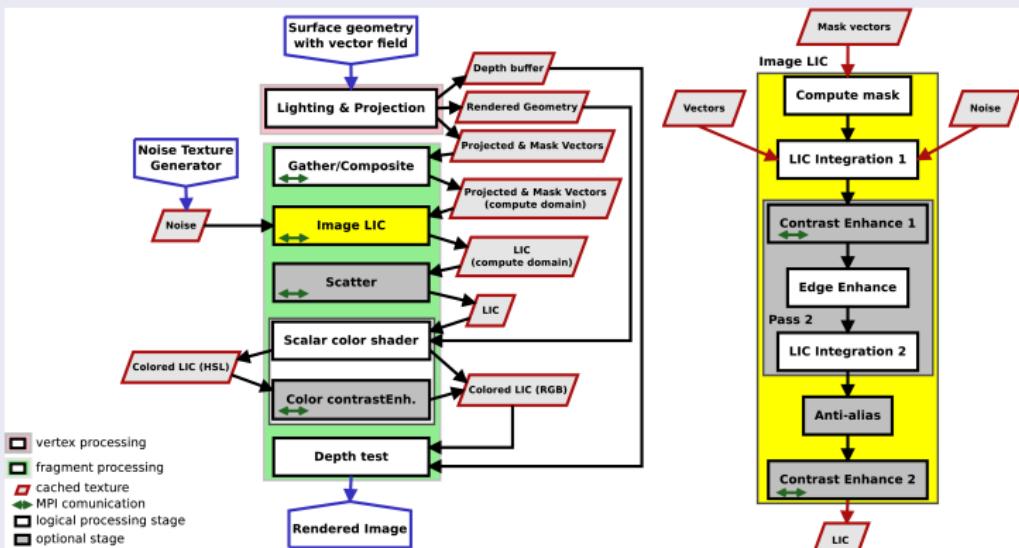
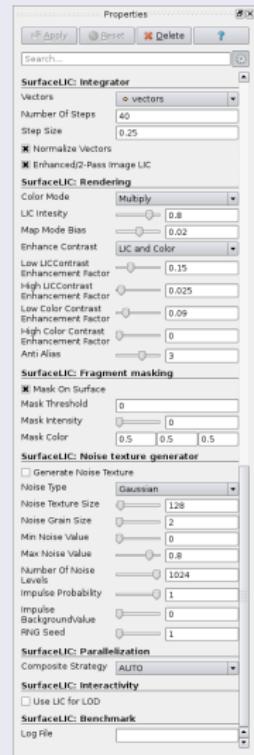
# Visual and perceptual issues

## Color Space Histogram Stretching: Applied on lightness channel after HSL conversion



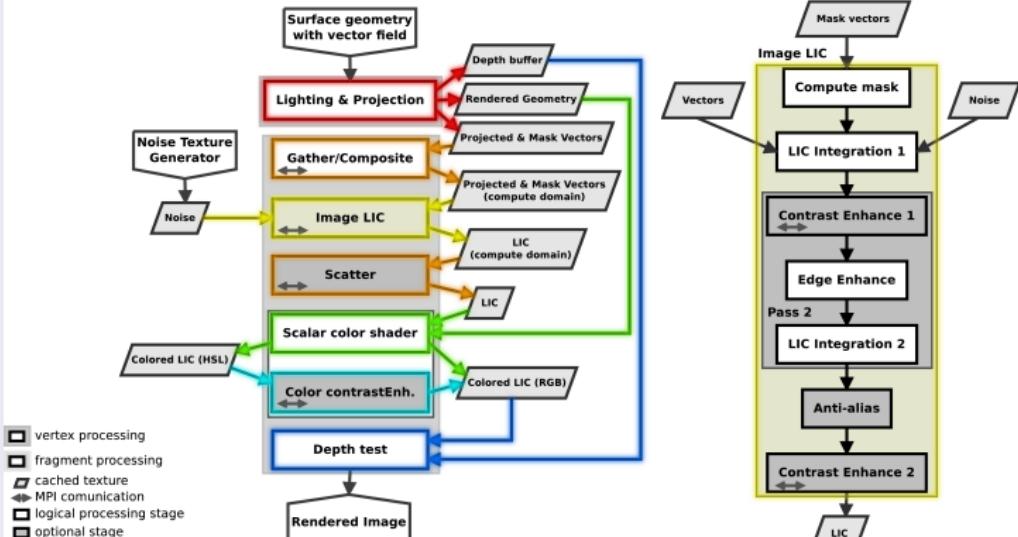
# Interactivity

efficiently respond to user tweaks



# Interactivity

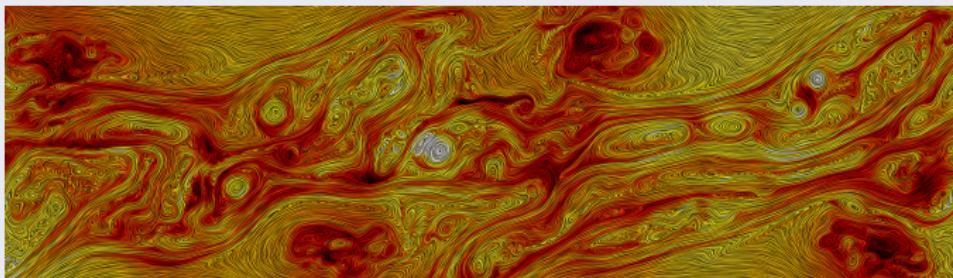
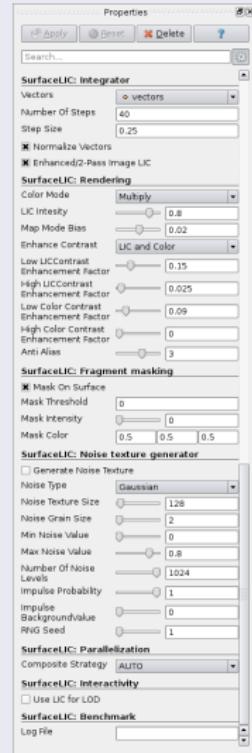
efficiently respond to user tweaks



- cache each stage of the pipeline in a texture
- group parameters by stage, only run modified and downstream stages

# Results

## What did we accomplish?



1. Apply image processing techniques to produce high contrast high dynamic range psuedocolored results ⇔ capture fine details and accurately represent the vector field
2. Expose many knobs ⇔ Choices are left to runtime in order to effectively deal with data, view, and surface variability
3. Optimizations for interaction ⇔ make interaction practical
4. Parallelizaiton ⇔ process very large datasets
5. Deployment in ParaView